


Memory Reach: GRU vs MLP under PPO on MiniGrid Memory

Anh Khoa Pham



Leibniz
Universität
Hannover

Code, studies and
winning checkpoints



1 Motivation

- **POMDP**: only a 7×7 view is seen
- **Memory** must live in a state h
- **Full BPTT** costs memory $\propto$ length
- **Truncation** cuts gradients to L

Core question: how far back must the gradient reach, and what does that reach cost?

250 runs · $\approx$ 250 GPU-hours · RTX 3060 12 GB · 16 CPUs


2 Problem Setting

The environment

- MiniGrid-MemoryS17Random-v0
- Cue (key or ball) in the west room
- Agent spawns facing away from it
- T-junction at the far east end
- Only the matching prong pays
- Reward $1 - 0.9 \cdot \text{steps} / \text{max_steps}$

The two encoders

- **MLP**: linear layers with ReLU, memory less
- **GRU**: carries h , with memory



Left: the maze. Right: everything the agent sees.

5 Key Insights


- 1 **Carrying h forward is not what matters.** Every GRU arm does it; only the reach of the gradient differs.
- 2 **What matters is the cue→decision gap.** Not the episode length. ≤ 8 steps in 62 % of mazes, ≤ 4 in 0%.
- 3 **The reach costs memory, not time.** Full BPTT is the fastest GRU arm (1.3 h); it pays in VRAM.

Recommendation. Full BPTT; $L = 8$ on a smaller GPU.


3 Approach

Network

actor-critic skeleton



GRU encoder, unrolled



PPO objective

$$\rho_t = \pi_\theta(a_t | s_t) / \pi_{\theta_{old}}(a_t | s_t)$$

probability ratio

$$\mathcal{L}^\pi = -\mathbb{E}_t[\min(\rho_t \hat{A}_t, \text{clip}(\rho_t, 1 - \epsilon, 1 + \epsilon) \hat{A}_t)]$$

clipped policy loss

$$\mathcal{L}^V = \mathbb{E}_t[(V_\theta(s_t) - \hat{R}_t)^2]$$

value loss

$$\mathcal{L} = \mathcal{L}^\pi + c_v \mathcal{L}^V - c_e \mathcal{H}[\pi_\theta]$$

what is backpropagated


Hyperparameter Optimization

- **Surrogate model:** Optuna TPE sampler, seeded at 42
- **Budget:** 50 trials per arm $\times$ 5 arms = 250 training runs
- **Seeds per trial:** 5, each 1000 iterations
- **Objective:** mean – std of return over the 5 seeds

4 Results

arm	return	success rate	steps	minibatch	time
MLP (memoryless)	0.556 ± 0.029	0.592 ± 0.027	24.3 ± 2.3	1024	1.3 h
GRU, L = 1	0.573 ± 0.064	0.588 ± 0.068	9.8 ± 4.1	4096	2.6 h
GRU, L = 4	0.492 ± 0.001	0.500 ± 0.000	7.0 ± 0.0	4096	1.6 h
GRU, L = 8	0.958 ± 0.004	1.000 ± 0.000	16.8 ± 1.4	4096	2.0 h
GRU, full BPTT	0.958 ± 0.002	1.000 ± 0.000	16.8 ± 0.9	512	1.3 h

Winning trial per study; mean $\pm$ std across the five seeds.



6 Limitations & Future Work

- One task, one architecture, one algorithm
- **GRU is slow:** the recurrence runs step by step
- Next: **LSTM** for a bigger environment
- Next: **transformer** for a faster training time

References. Hausknecht & Stone 2015 · Kapturowski et al. 2019 (R2D2) · Ni et al. 2022 · Schulman et al. 2017 (PPO) · Chevalier-Boisvert et al. 2023 (MiniGrid)